

Embedding Agentic Traces

A trace representation should be faithful to content and invariant to authorship

Helivan Research — working summary, August 2026

The problem

Modern generative systems produce logs of their actions and decisions—referred to as an *agentic trace*—while processing information. Traces often implicitly or explicitly encode information such as the particular type of tool that was used, the harness, etc. that may or may not be useful for determining if the trace produced a useful output. Further, the signal from using an off-the-shelf embedding function to encode trace similarity is often dominated by information that is potentially irrelevant for correctness, such as the harness or model used: in our corpus, nearest neighbors of raw trace embeddings are same-system 68–83% of the time, same-scaffold at $5\times$ chance, and the exact agent configuration is retrievable at 0.82 vs. 0.43 chance even among same-model, same-harness siblings—while agreement with outcome sits at chance (AUC 0.52).

The goal of this work is simple: **develop a method for embedding agentic traces that better captures similarity of the work performed.** In developing such a method, we have identified six desirable properties of an agentic trace embedding, organized into three groups:

Faithful to content — what happened should drive similarity.

- **Task fidelity:** traces of different systems working the same problem should be similar. *Measured by* cross-system same-task nearest-neighbor retrieval.
- **Behavioral fidelity:** among traces of the same problem, similar approaches and results should be similar. *Measured by* within-task coherence of same-outcome traces (AUC of same- vs. different-outcome distances).

Invariant to authorship — who produced it should not.

- **Harness invariance:** scaffold formatting and tool idiom should not drive similarity. *Measured by* within-task same-scaffold retrieval vs. chance.
- **Model-family invariance:** the vendor/lineage of the underlying LLM should not drive similarity. *Measured by* within-task same-family retrieval vs. chance.
- **Identity invariance:** the exact system configuration (version, prompts, settings) should not be recoverable. *Measured by* exact-submission retrieval among same-model, same-harness siblings vs. chance.

Reliable — similarity should be a measurement, not noise.

- **Stability:** the embedding should be stable under agent stochasticity, judge resampling, and embedder choice. *Measured by* replicate-run and judge-resample agreement, per trace and aggregated over instances.

qubric

Our method for embedding agentic traces is based on a query-specific rubric. In particular, for each task instance an LLM (Model 1) reads the public problem statement and writes a query-specific rubric comprised of 1) understanding, 2) localization, 3) reproduction, 4) editing, 5) verification, and 6) final state. Given a trace corresponding to the task, a judge LLM (Model 2) extracts a short factual description per section. The short descriptions are then embedded with an off-the-shelf embedding model, centered against the cross-system median per instance, and concatenated. This approach enforces the equivalence relations the space should respect by respecting the rubric and using a fixed judge LLM as the extractor.

Evidence

We evaluate qubric on 107 system configurations from the SWE-bench Verified leaderboard (a 500-task benchmark; qubric judging is on a 20-task panel, with the 150-task scale-up in progress). We use gpt-5.4-mini for Model 1 and gpt-5.4-mini for Model 2 in the primary configuration; robustness to Model 2 is reported below. We compare qubric along the six axes described above against four baselines:

- **raw trace:** embed the first/last 8K tokens of the rendered trace directly—the off-the-shelf default;
- **free-form summary:** Model 2 summarizes the trace with no rubric—tests whether structure matters;
- **generic rubric:** the same six sections, fixed across instances—tests whether query-specificity matters;
- **verdict questions:** the rubric phrased as evaluative questions with short answers—tests descriptions vs. verdicts;

across **seven off-the-shelf embedding functions** spanning three families and two orders of magnitude of size. Per-pillar results for the baselines are shown in Fig. 2.

Qubric produces the same profile shift under all seven embedders (Fig. 1): *task fidelity* is retained (~ 0.95) and *behavioral fidelity* retained or improved, while every authorship axis collapses to $\approx$ chance—*harness* $0.31 \rightarrow 0.08$, *model family* $0.20 \rightarrow 0.12$, *identity* $0.83 \rightarrow 0.09$. The properties are separately measurable (double dissociations across six representations), so these are genuinely different axes being moved, not one metric relabeled.

The instance-conditioning is the active ingredient (Fig. 4): the construction ranking is strict under both readouts (qubric > verdict questions > generic rubric > free-form summary > raw), reproduces under a local embedder, and no single




Figure 1: Sensitivity radars across seven embedders. Outside = desirable. Raw embeddings (navy) are authorship-dominated; qubric (amber) collapses Identity / Model family / Harness to chance while holding Task and Behavior.




Figure 2: Baselines by pillar: each cell is the raw metric for that axis (columns) under that representation (rows); color is desirability (blue = ideal). Raw-trace variants (head/tail/centered) read authorship; qubric is the only row that is simultaneously task-faithful and authorship-invariant.

rubric section carries the effect (best single section 0.091 vs. 0.078 for the union). Judge choice degrades gracefully: an open frontier judge (DeepSeek-v3.1, 0.083/0.079) matches the closed reference (gpt-5.4-mini, 0.078/0.078) within panel resolution, and a weak judge *with* a qubric outperforms a strong judge without one.

Stability: two judges agree on a single trace only 12–28% of the time, but aggregating ~ 20 instances drives system-level agreement to 0.82–0.95; replicate runs of the same agent retrieve each other at $6\times$ chance. Per-trace noisy, per-agent reliable—and per-agent is what evaluation uses.

Payoff: query-efficient benchmarking

Qubric representations enter the DKPS/PKPS perspective space; a new agent’s full-benchmark resolve rate is predicted from a few probe traces against references’ cached coverage,

blended with the sample score. Evaluation is leave-one-LLM-out throughout (naive CV inflates accuracy $\sim 25\%$ via model-family leakage).

probes m	sample score	qubric pipeline	+selection
1	0.44	0.095	—
5	0.16	0.074	0.057
20	0.06	0.048–0.053	0.053

Table 1: MAE predicting true resolve rates (107 systems). One probe trace ≈ 8 –10 scored runs.

One trace carries far more than one bit: the strongest correctness-only competitor (CAPA-style error overlap) needs 15–20 *scored* probes to match one *trace* (0.37 MAE at $m=1$), and trails the ensemble at every budget. The reference pool, not probe count, is the binding resource: error falls $0.145 \rightarrow 0.078$ as the pool grows $5 \rightarrow 96$, still declining (Fig. 3). Every layer is ablated; CV-inside-selection and PCA-composed-with-selection are documented overfitting traps.

Caveats: panel differences < 0.01 are near resolution; one benchmark so far. Full log: FINDINGS.md (F1–F24), agentic-traces branch.




Figure 3: QUENCH: error predicting true resolve rates vs. probe budget m (left) and vs. reference-pool size n (right); leave-one-LLM-out, 107 systems.

Judge sensitivity: prediction error by construction and extractor
Cells: leave-one-LLM-out kNN MAE (resolve rate), 107 systems x 20 instances; navy = lower error.

gpt-oss-20b (open)	pending	pending	pending	0.113 ridge 0.100
gpt-4o-mini	0.110 ridge 0.089	pending	pending	pending
gpt-oss-120b (open)	pending	pending	pending	0.086 ridge 0.092
deepseek-v3.1 (open)	pending	pending	pending	0.083 ridge 0.079
gpt-5.4-nano	pending	pending	pending	0.083 ridge 0.081
gpt-5.4-mini	0.094 ridge 0.083	0.088 ridge 0.082	0.086 ridge 0.090	0.078 ridge 0.078
	Free-form summary	Generic rubric	Verdict questions	Qubric

Rubric writer fixed at gpt-5.4-mini (only cached writer); all cells embedded with nomic-embed-text-v1.5.
Naïve trace baseline (no judge): kNN 0.102.

Figure 4: Judge robustness: leave-one-LLM-out kNN MAE by construction and extraction judge (Model 2); navy = lower error. Query-specificity beats every alternative under the same judge, and a weaker judge with a qubric beats a stronger judge without one.